

Deep Learning and Temporal Data Processing

3 - Recurrent Neural Networks

Andrea Palazzi

June 21th, 2017

University of Modena and Reggio Emilia

Introduction

Credits

Introduction

test [1]

In **feedforward neural network** computation flows directly from input x through intermediate layers h to output y .

Conversely, some networks topology feature feedback connections, in other words model outputs are fed back into the model itself.

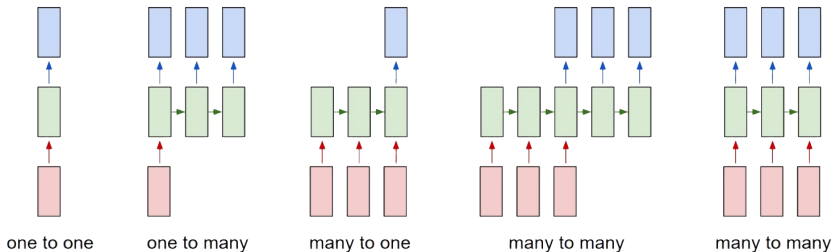
The term **recurrent neural networks** defines this family of models.


Recurrent neural networks (RNN) are **specialized for processing sequences**.

Similarly, we saw that convolutional neural networks feature specialized architecture for processing images.

RNNs boast a **much wider API with respect to feedforward neural networks**.

Indeed, these models can deal with *sequences* in the input, in the output or even both.






The vanilla RNN is provided with three sets of parameters:

- U maps inputs to the hidden state
- W parametrizes hidden state transition
- V maps hidden state to output

System dynamics is as simple as:

$$\begin{cases} \mathbf{h}^{(t)} = \phi(\mathbf{W} \mathbf{h}^{(t-1)} + \mathbf{U} \mathbf{x}^{(t)}) \\ \mathbf{o}^{(t)} = \mathbf{V} \mathbf{h}^{(t)} \end{cases} \quad (1)$$




The hidden state $\mathbf{h}^{(t)}$ can be intuitively viewed as a *lossy* summary of the sequence of past inputs fed to the network, in which are stored the main task-relevant aspects of the past sequence of inputs up to time t .

Since the an input sequence of arbitrary length $(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(t)})$ is mapped into a fixed size vector $\mathbf{h}^{(t)}$, this summary is necessarily lossy.

A recurrent computational graph can be unfolded into a sequential computational graph with a repetitive structure.

$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)}; \theta)$$



how to unroll a recursive graph

Vanishing and exploding gradient problem.

Vanishing and exploding gradient problem.

No matter how's the network topology, during backpropagation the network is unfolded in a DAG, so there are no loops.

Credits

These slides heavily borrow from the following Stanford course:

- <http://cs231n.stanford.edu/>

if you want to deepen your knowledge of these concepts, I'd really suggest you to start from here!

Also, nice convolution animations are taken from here:

- https://github.com/vdumoulin/conv_arithmetic

[1] G. Cybenko.

Approximation by superpositions of a sigmoidal function.

Mathematics of Control, Signals, and Systems (MCSS), 2(4):303–314, 1989.